

# Lab 8

## Decision Trees



Escola Tècnica Superior d'Enginyeria  
de Telecomunicació de Barcelona

UNIVERSITAT POLITÈCNICA DE CATALUNYA

# Introduction

## Objectives

- Learn how to use decision trees to solve a classification problem
- Learn to prune a decision tree
- Learn to train a random forest classifier
- Solve the Ionosphere classification problem with decision trees and random forests
- Solve the MNIST classification problem with random forests.

# Decision Trees in sklearn

Class: `DecisionTreeClassifier`

Scikit-Learn uses the **CART algorithm**, which produces only binary trees. CART constructs binary trees using the feature and threshold that yield the largest information gain at each node.

The main parameters are:

**criterion:** The function to measure the quality of a split. Supported criteria are “gini” (default) for the Gini impurity and “entropy” for the information gain.

**max\_depth:** The maximum depth of the tree. If None (default), then nodes are expanded until all leaves are pure or until all leaves contain less than **min\_samples\_split** samples.

**min\_samples\_split:** The minimum number of samples required to split an internal node (default=2)

**min\_samples\_leaf** The minimum number of samples required to be at a leaf node (default 1). A split point at any depth will only be considered if it leaves at least `min_samples_leaf` training samples in each of the left and right branches. This may have the effect of smoothing the model, especially in regression.

**class\_weight** Weights associated with classes (default None).

# Decision Trees in sklearn

```
clf = sklearn.tree.DecisionTreeClassifier(criterion='entropy' )  
clf.fit(X, y)
```

## #Plot the tree

```
sklearn.tree.plot_tree(clf, ax=fig.subplots());
```

## #Export in textual format

```
t = export_text(clf)
```

## #Check parameters used

```
clf.get_params()
```

## #Feature importances

```
clf.feature_importances_
```

## #Parameters for regularization

```
max_depth
```

```
min_samples_split
```

```
min_samples_leaf
```


```
min_weight_fraction_leaf
```

```
max_leaf_nodes
```

```
max_features
```

# Bias and variance

Decision trees are large variance predictors: a slight modification in the dataset can dramatically change the decision boundaries.



# Random Forests in sklearn

To improve the predictive accuracy and control over-fitting, multiple decision tree classifiers can be trained and later combined.

In random forests, randomness is introduced in the classifier construction by :

1. Each tree is built from a **sample drawn with replacement from the training set**.
2. When splitting each node during the construction of a tree, the best split is found from a **random subset from the input features**.

Random forests achieve a reduced variance by combining diverse trees, sometimes at the cost of a slight increase in bias. In practice the variance reduction is often significant hence yielding an overall better model

Random foeste class: RandomForestClassifier

```
clf =  
sklearn.ensemble.RandomForestClassifier(criterion='entropy',  
min_samples_leaf=5)  
clf.fit(X, y)
```

# Random forest

One of the most powerful techniques for shallow learning to date.

We generate an ensembles of trees by

1. Using **bagging** plus
2. On each node we select a **random subset** of  $m$  features to decide the best. The subset changes randomly from each node to the next.

The value of  $m$  and the number of trees  $M$  are hyperparameter.

We can keep some of the **explainability** of decisions trees by checking how many times over all trees some features are selected weighted by the number of samples in the node.

# Bagging

The ensemble of classifiers is trained using different subsets of data, possibly with replacement. It reduces variance.

